

**RSAC** | 2025  
Conference

Many Voices.  
**One Community.**

SESSION ID: CLS-M01

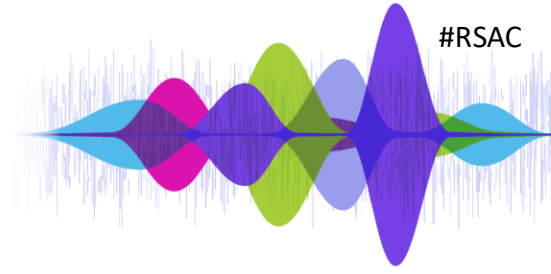
# Securing Cloud Access with Kubernetes Workload Identity

**Eric Johnson**

Principal Security Engineer, Puma Security  
Senior Instructor, SANS Institute  
[linkedin.com/in/eric-m-johnson/](https://linkedin.com/in/eric-m-johnson/)



# Disclaimer



Presentations are intended for educational purposes only and do not replace independent professional judgment. Statements of fact and opinions expressed are those of the presenters individually and, unless expressly stated to the contrary, are not the opinion or position of RSA Conference LLC or any other co-sponsors. RSA Conference LLC does not endorse or approve, and assumes no responsibility for, the content, accuracy or completeness of the information presented.

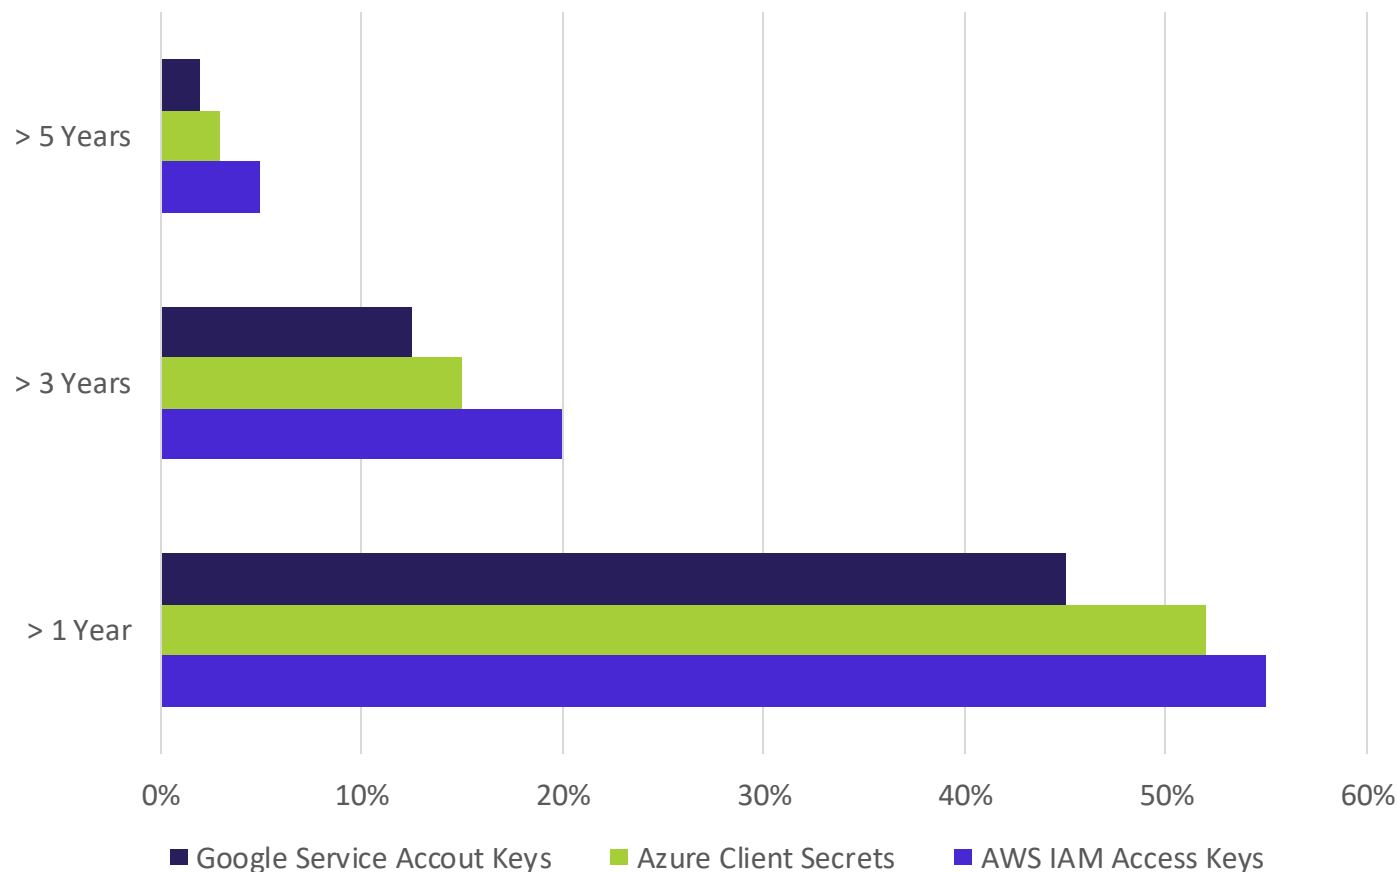
Attendees should note that sessions may be audio- or video-recorded and may be published in various media, including print, audio and video formats without further notice. The presentation template and any media capture are subject to copyright protection.

© 2025 RSA Conference LLC or its affiliates. The RSAC and RSAC CONFERENCE logos and other trademarks are proprietary. All rights reserved.

# Cloud Attack Techniques

- Static credentials primarily support programmatic scripts and applications
- Rotating static credentials is challenging, so we avoid it
- Lost or stolen static credentials are the initial access method in ~ 66% of cloud breaches
- Sources:
  - AWS re:Inforce TDR432: New tactics and techniques for proactive threat detection
  - 2024 Datadog State of Cloud Security

Static Credential Age By Cloud Provider



# Workload Identity GitHub Repository

Nymeria is an open-source repository with many cross-cloud static credential and workload identity scenarios:

- Kubernetes
  - AWS EKS, Azure AKS, GKE
- Serverless functions
  - AWS Lambda, Azure & Google Cloud Functions
- GitHub Actions / Azure virtual machines

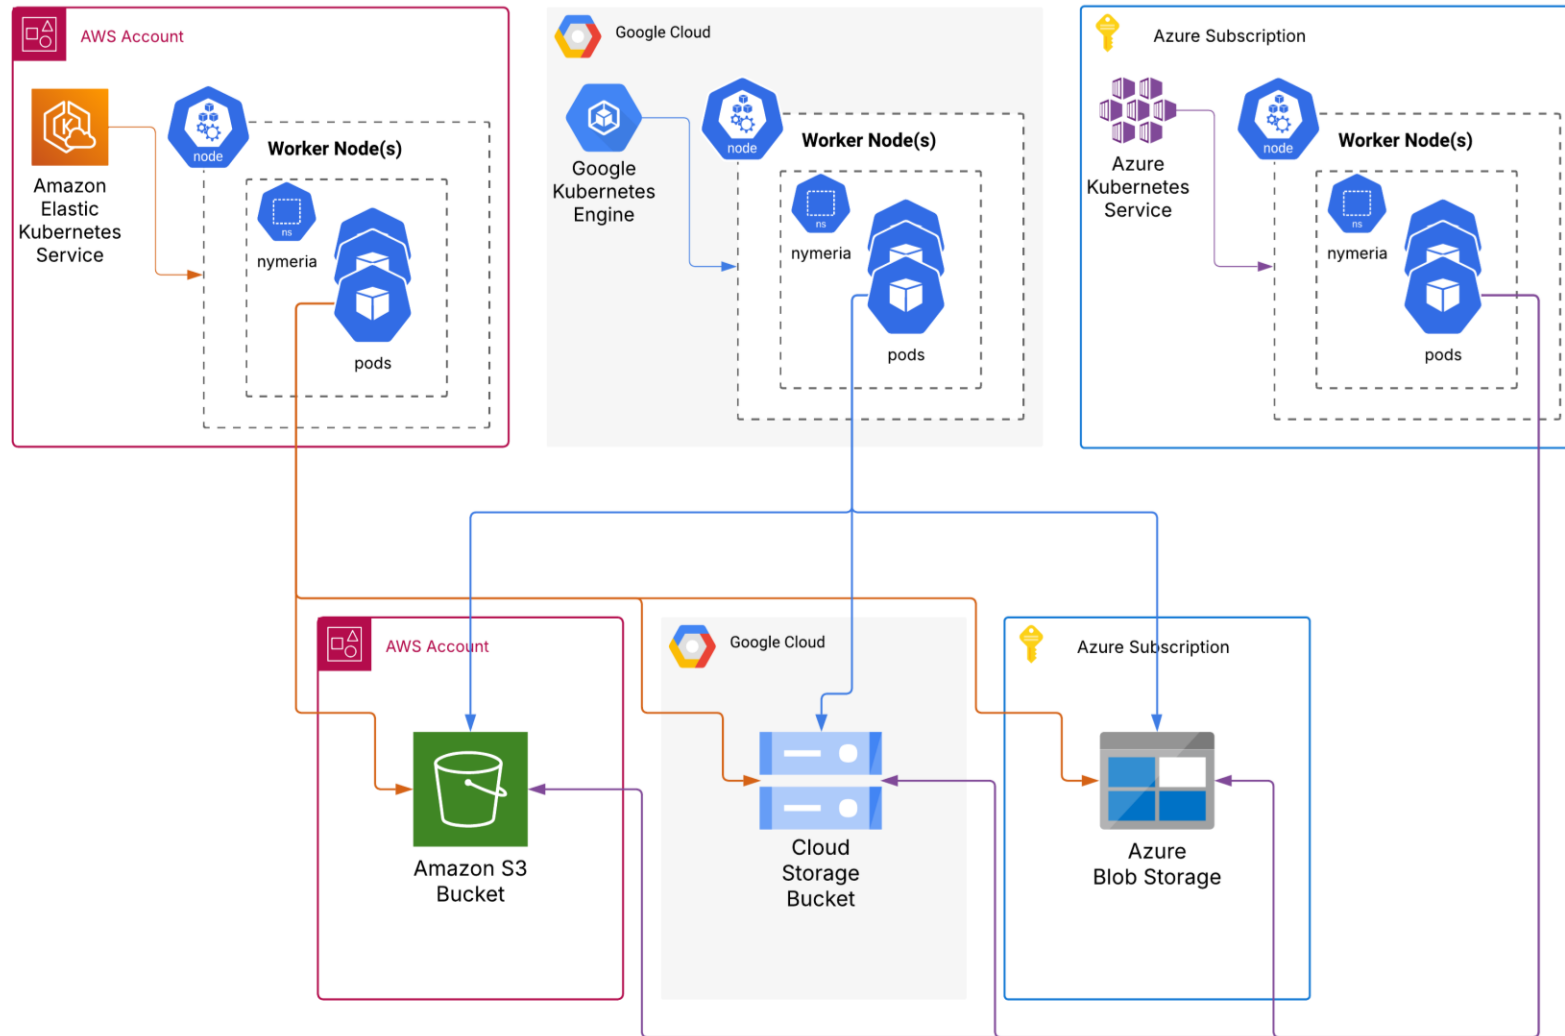


**NYMERIA**



<https://github.com/pumasecurity/nymeria>

# The Problem: Kubernetes Cross Cloud Authentication



# Session Goals

- Review Kubernetes pod credential anti-patterns
- Learn how to enable Kubernetes workload identity
- Configure pod service account identity tokens for intra-cloud access to resources
- Extend pod service account identity tokens to meet external cloud provider audience and expiration requirements
- Programmatically authenticate from a pod to AWS, Azure, and Google Cloud APIs

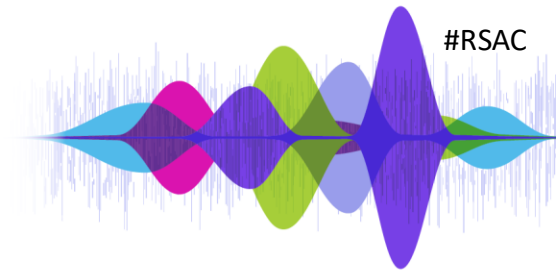


# Kubernetes Pod Credential Anti-Patterns

Securing Cloud Access with Kubernetes  
Workload Identity

Many Voices.  
**One Community.**

# Cloud Provider Static Credential Types



Each cloud provider's Identity & Access Management (IAM) service provides an option for creating static, long-lived credentials:



**IAM User  
Access Keys**



**Service Principal  
Client Secrets**

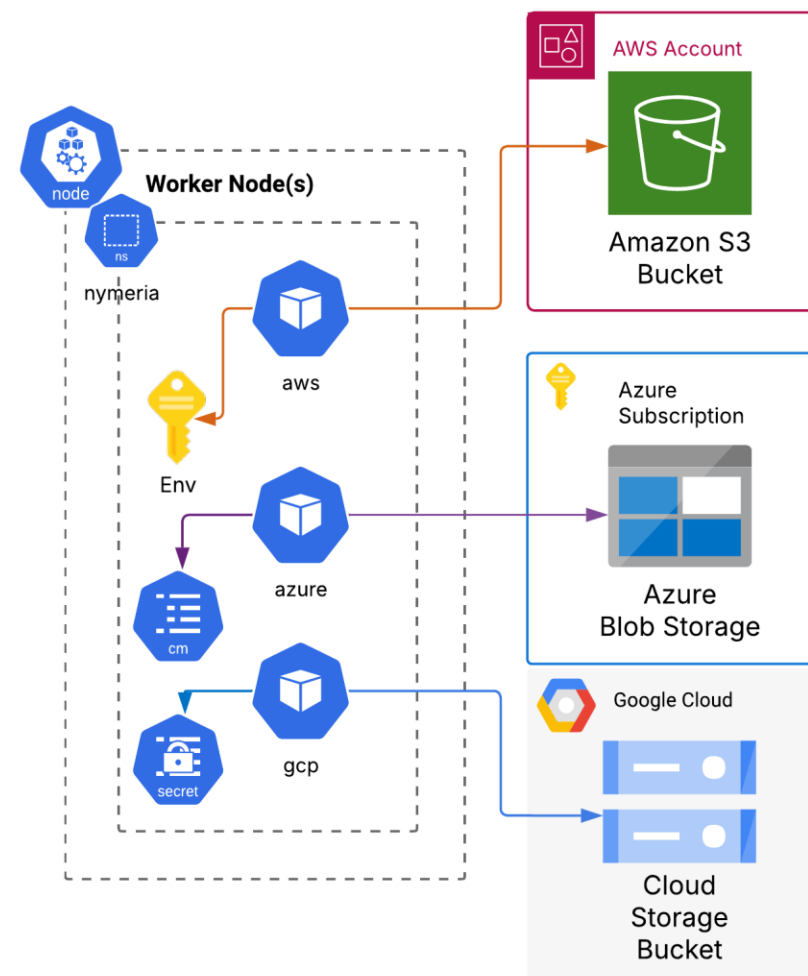


**Google Service  
Account Keys**



# Pod Credential Anti-Patterns

- Generating static, long lived credentials for accessing cloud provider APIs
- Storing static cloud credentials in Kubernetes objects:
  - ConfigMap
  - Secrets
  - Annotations
- Provisioning credentials to pods through environment variables or file system volume mounts

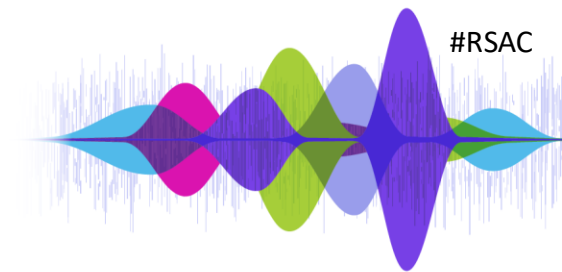


# Kubernetes Pod ConfigMap Mount

Discovering static credentials provisioned using a *ConfigMap*:

```
1  $ kubectl describe pods -n static-credential -l cloud=gcp
2
3  Name:          nymeria-gcloud-8478f4cb8f-bs8kr
4  Namespace:     static-credential
5  ...
6  Containers:
7    gcloud:
8      ...
9      Image: gcr.io/google.com/cloudsdktool/google-cloud-cli:latest
10     Mounts:
11       /mnt/gcp/sa from gcp-sa-key (rw)
12   ...
13  Volumes:
14     gcp-sa-key:
15       Type:          ConfigMap (a volume populated by a ConfigMap)
16       Name:          gcp-sa-key
```

# Kubernetes ConfigMap Exfiltration



Exfiltrating *ConfigMap* data containing static credentials:

```
1  $ kubectl describe configmaps -n static-credential gcp-sa-key
2
3  Name:          gcp-sa-key
4  Namespace:     static-credential
5
6  Data
7  ====
8  credentials.json:
9  ----
10 {
11     "type": "service_account",
12     "project_id": "your-project",
13     "private_key": "-----BEGIN PRIVATE KEY..."
14     "client_email": "sa@your-project"
15     ...
16 }
```

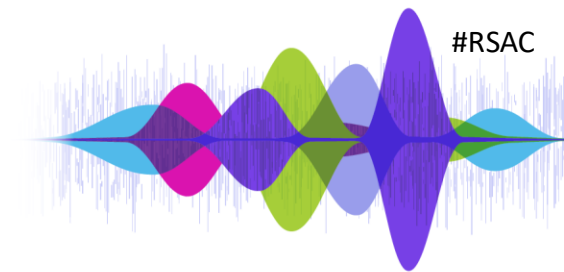
# Kubernetes Pod Environment Variables



Discovering static credentials provisioned in a pod's environment variables:

```
1 $ kubectl describe pods -n static-credential -l cloud=aws
2
3 Name:          nymeria-aws-58fcfc8969-t7qwr
4 Namespace:     static-credential
5 Containers:
6   awscli:
7     Container ID:  containerd://...
8     Image:         amazon/aws-cli:latest
9     ...
10  Environment:
11    AWS_ACCESS_KEY_ID: <set to the key
12    'aws_secret_access_key_id' in secret 'aws-iam-user'>
13    AWS_SECRET_ACCESS_KEY: <set to the key
14    'aws_secret_access_key' in secret 'aws-iam-user'>
```

# Kubernetes Secrets Exfiltration



Exfiltrating secret values containing static credentials:

```
1 $ kubectl get secrets -n static-credential -o json aws-iam-user
2
3 {
4   "apiVersion": "v1",
5   "data": {
6     "aws_secret_access_key":
7       "SElacHZmS1lCU3M3cWNyZ3ZZYURUSnJMc nE0S2l1MGQ5aDFYcG5MTg==",
8     "aws_secret_access_key_id": "QUtJQVNaWTJaU1U2WVVFWE5HSTY="
9   },
10  "kind": "Secret",
11  ...
12  },
13  "type": "Opaque"
14 }
```

# Kubernetes Credentials Mitigations

Common mitigations....

- Version control secrets scanning



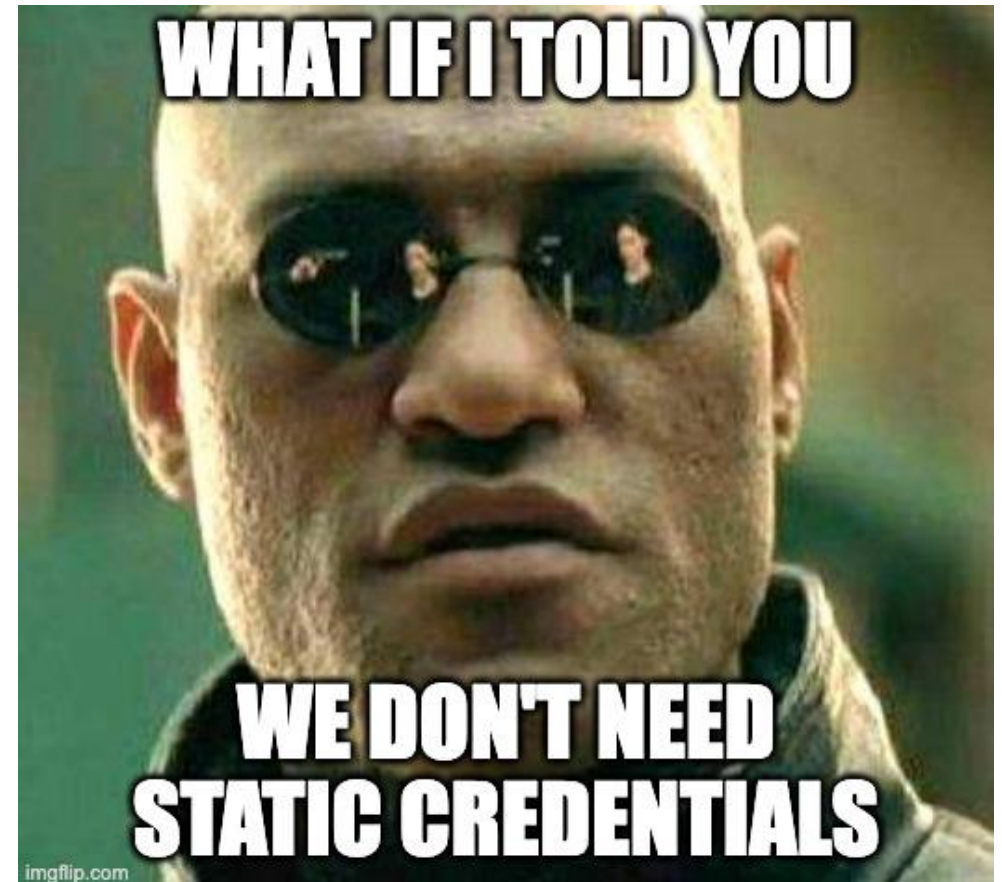
- Kubernetes secrets management integrations



- Compromised credential detection



But....



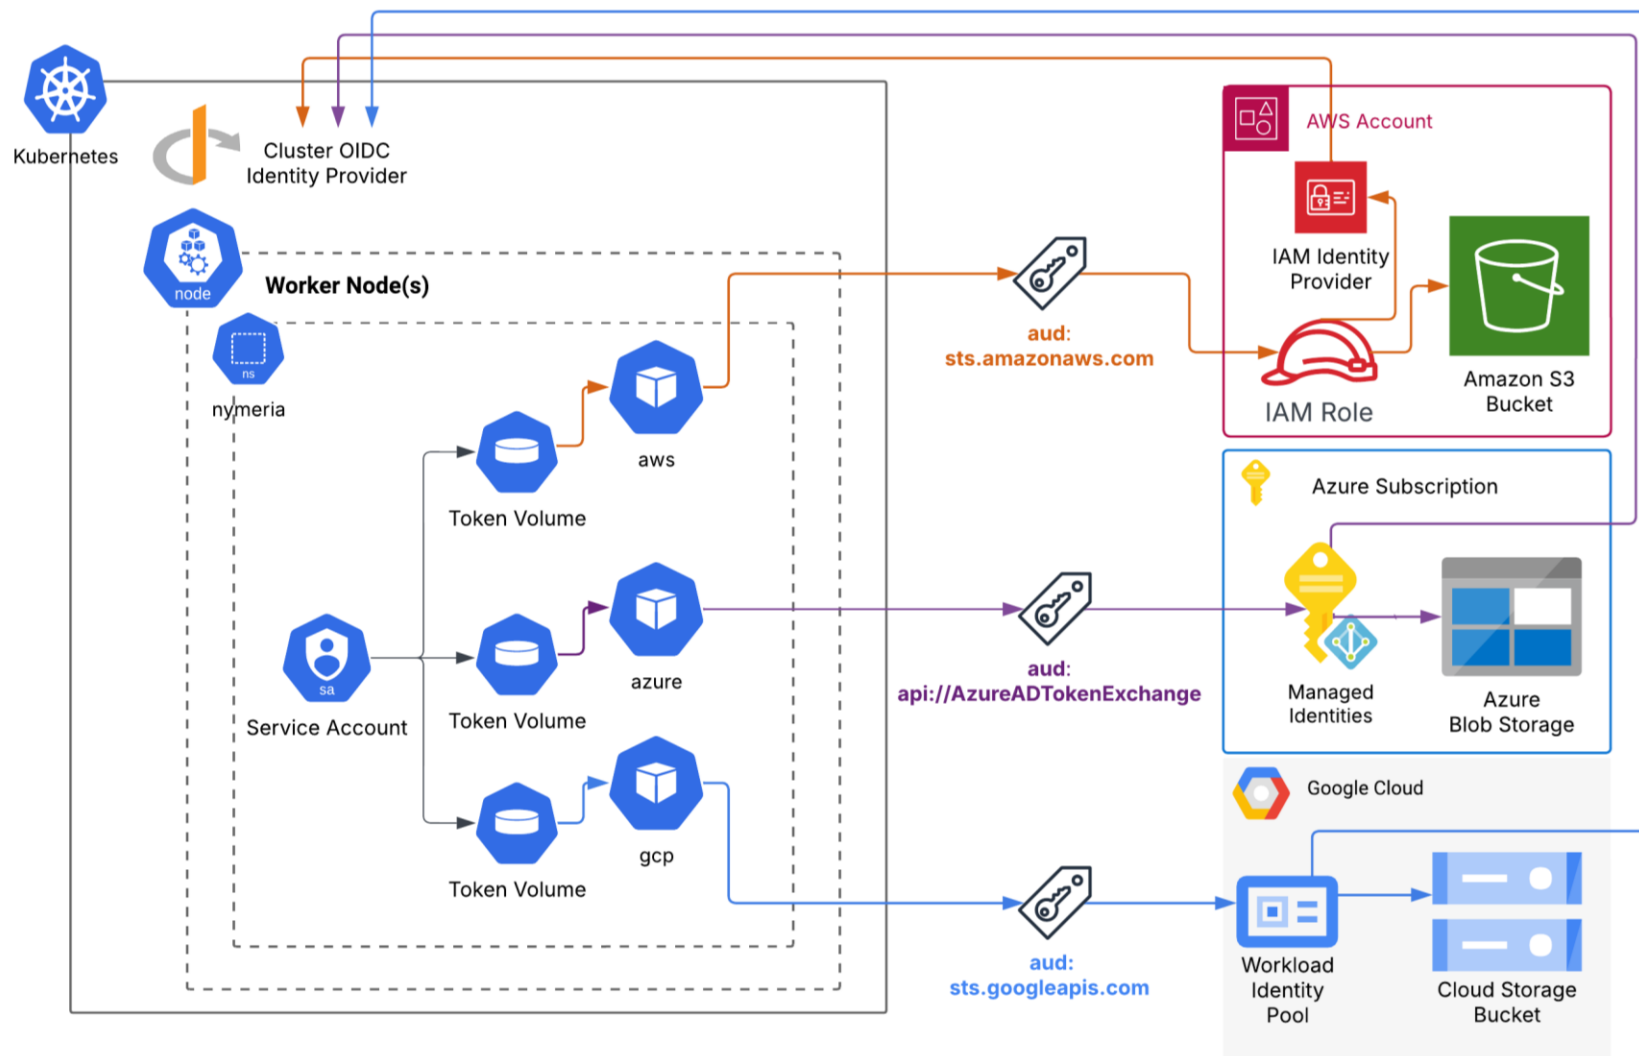


# Kubernetes Workload Identity Federation

Securing Cloud Access with Kubernetes  
Workload Identity

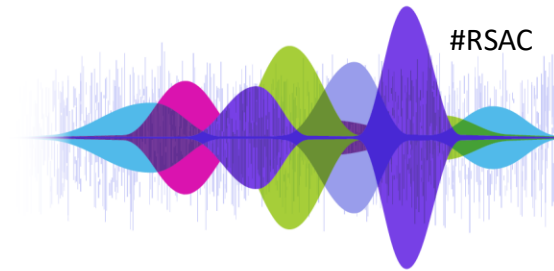
Many Voices.  
**One Community.**

# Kubernetes Workload Identity Federation





# Cloud Kubernetes Service Capabilities



Cloud managed Kubernetes service add-ons enable workload identity for pods accessing the cloud provider APIs:

AWS Elastic Kubernetes Service (EKS)



Create an IAM Roles for Service Accounts (IRSA) OIDC identity provider or install the EKS Pod Identity add-on

Google Kubernetes Engine (GKE)



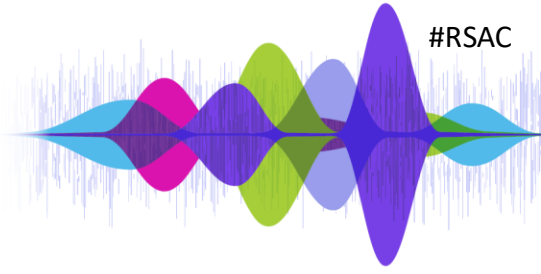
Enable workload identity and create a namespace for Kubernetes service accounts (automatically enabled for auto pilot)

Azure Kubernetes Service (AKS)



Enable the workload identity and OIDC issuer add-on options for the cluster

# Kubernetes Workload Identity Configuration



Kubernetes workload identity allows pods to securely access cloud provider APIs without managing static credentials:

1. Create a Kubernetes service account
2. Set required service account annotations (varies by cloud provider)
3. Assign the service account to the pod
4. Observe the service account identity tokens are automatically fed into the pod and rotated by the cluster



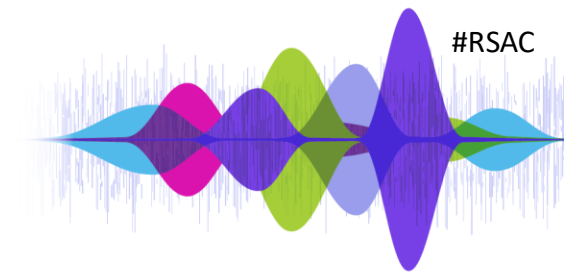
# GKE / Google Cloud Workload Identity



- Kubernetes *ServiceAccount* resources provide a unique identity for workloads running inside a namespace
- Deployments (or pods) using the *serviceAccountName* will have a signed OIDC identity token projected into a volume mount
- GKE specific service account annotations are ***not*** required

```
1  ---
2  apiVersion: v1
3  kind: ServiceAccount
4  metadata:
5    name: "nymeria"
6    namespace: workload-identity
7  ---
8  apiVersion: apps/v1
9  kind: Deployment
10 ...
11 spec:
12   ...
13   template:
14     ...
15     spec:
16       serviceAccountName: nymeria
17       containers:
18         - name: gcloud
```

# GKE Workload Identity Tokens



Service account identity tokens are automatically mounted into a volume for the pod to use for intra-cluster and Google cloud API access:

```
1 $ kubectl exec -n workload-identity -it $(kubectl get pod -n workload-  
2 identity -l cloud=gcp -o json | jq -r '.items[0].metadata.name') --  
3 cat /var/run/secrets/kubernetes.io/serviceaccount/token
```

Decoding the service account identity token reveals the GKE cluster's default issuer and audience, as well as the node, pod, and service account:

```
1 {  
2   "aud": [ "https://container.googleapis.com/v1/projects/my-project  
3           /locations/my-region/clusters/nymeria" ],  
4   "iss": "https://container.googleapis.com/v1/projects/my-project  
5           /locations/my-region/clusters/nymeria",  
6   "sub": "system:serviceaccount:workload-identity:nymeria"  
7   "kubernetes.io": {...} ...}
```



# GKE Workload Identity Policy Binding

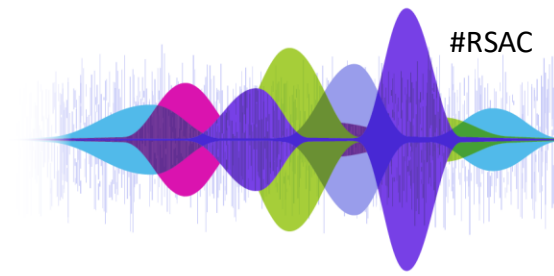
GKE service accounts can be directly added to project / resource level IAM permission bindings:

- GKE automatically creates a workload identity pool in the project
- Example: granting a GKE service account principal the storage object viewer role on a bucket:

```
1  $ gcloud storage buckets get-iam-policy gs://nymeria-123456
2
3  bindings:
4    role: roles/storage.objectViewer
5    - members:
6      - principal://iam.googleapis.com/projects/#####/locations/
7        global/workloadIdentityPools/my-project.svc.id.goog/subject
8          /ns/workload-identity/sa/nymeria
```



# AWS EKS IRSA / AKS Workload Identity



EKS *ServiceAccount* resources use annotations to configure workload identity:



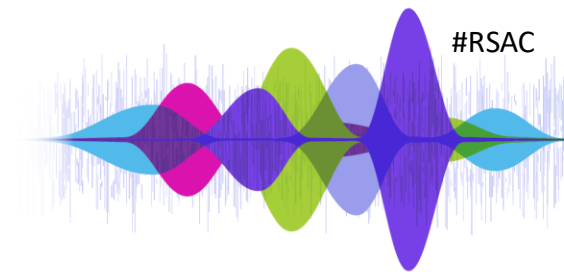
```
1  apiVersion: v1
2  kind: ServiceAccount
3  metadata:
4    name: "nymeria"
5    namespace: workload-identity
6    annotations:
7      eks.amazonaws.com/role-arn:
8        arn:aws:iam::0123456789012:
9        role/nymeria
10     eks.amazonaws.com/token-
11     expiration: "3600"
```

AKS *ServiceAccount* resources use annotations to configure workload identity:



```
1  apiVersion: v1
2  kind: ServiceAccount
3  metadata:
4    name: "nymeria"
5    namespace: workload-identity
6    annotations:
7      azure.workload.identity/
8      client-id: your-client-id
9      azure.workload.identity/
10     tenant-id: your-tenant-id
11     azure.workload.identity/
12     service-account-token-
13     expiration: "3600"
```

# AWS EKS IRSA Identity Tokens



Service account identity tokens are automatically mounted into a volume for the pod to use for AWS API authentication:

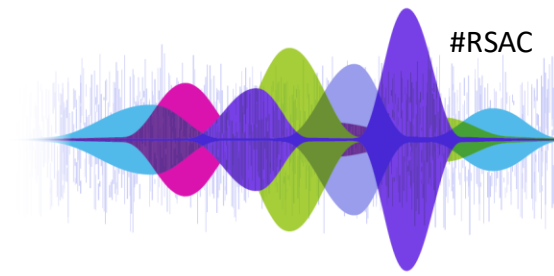
```
1 $ kubectl exec -n workload-identity -it $(kubectl get pod -n workload-  
2 identity -l cloud=aws -o json | jq -r '.items[0].metadata.name') --  
3 cat /var/run/secrets/eks.amazonaws.com/serviceaccount/token
```

Decoding the service account identity token reveals the EKS cluster's default issuer and audience, as well as the node, pod, and service account:

```
1 {  
2   "aud": [ "sts.amazonaws.com" ],  
3   "iss": "https://oidc.eks.us-west-1.amazonaws.com/id/my-cluster-id",  
4   "sub": "system:serviceaccount:workload-identity:nymeria"  
5   "kubernetes.io": {...}  
6   ...  
7 }
```



# AWS EKS IRSA Role Assumption




1. AWS Identity Provider resource trusting identity tokens issued from the EKS cluster OIDC identity provider:

oidc.eks.us-west-1.amazonaws.com/id/my-cluster-id [Info](#)

[Assign role](#) [Delete](#)

**Summary**

<b>Provider</b> oidc.eks.us-west-1.amazonaws.com/id/my-cluster-id	<b>Provider Type</b> OpenID Connect
<b>Creation Time</b> February 22, 2025, 16:46 (UTC-06:00)	<b>ARN</b>  <code>arn:aws:iam::123456789012:oidc-provider/oidc.eks.us-west-1.amazonaws.com/id/my-cluster-id</code>

[Audiences \(1\)](#) [Endpoint verification](#) [Tags](#)

**Audiences (1)** [Actions](#)

Also known as client ID, audience is a value that identifies the application that is registered with an OpenID Connect provider.

< 1 >

Audience
<input type="radio"/> sts.amazonaws.com

2. AWS IAM Role trust policy granting the EKS pod's service account permissions to assume the role:

```
1  { "Effect": "Allow",
2    "Principal": {
3      "Federated":
4        "arn:aws:iam::123456789012:
5          oidc-provider/oidc.eks.us-west-1
6            .amazonaws.com/id/my-cluster-id"
7    },
8    "Action":
9      "sts:AssumeRoleWithWebIdentity",
10   "Condition": {
11     "StringEquals": {
12       "oidc.eks.us-west-1.amazonaws.com
13         /id/my-cluster-id:sub":
14       "system:serviceaccount:workload-
15         identity:nymeria"
16     }
17   }
```





# Azure Kubernetes Workload Identity Tokens

Service account identity tokens are automatically mounted into a volume for the pod to use for Azure API authentication:

```
1 $ kubectl exec -n workload-identity -it $(kubectl get pod -n workload-  
2 identity -l cloud=azure -o json | jq -r '.items[0].metadata.name') --  
3 cat /var/run/secrets/azure/tokens/azure-identity-token
```

Decoding the service account identity token reveals the AKS cluster's default issuer and audience, as well as the node, pod, and service account:

```
1 {  
2   "aud": [ "api://AzureADTokenExchange" ],  
3   "iss": "https://eastus2.oic.prod-aks.azure.com/tenant-id/cluster-id/",  
4   "sub": "system:serviceaccount:workload-identity:nymeria"  
5   "kubernetes.io": {...}  
6   ...  
7 }
```



# Azure Managed Identity Federation

1. Azure managed identities support logging in using a federated credential (OIDC JWT)
2. Each federated credential requires the cluster issuer, audience, and service account

nymeria-49d0fec8 | Federated credentials

Managed Identity

Search

- Overview
- Activity log
- Access control (IAM)
- Tags
- Azure role assignments
- Associated resources (preview)
- Resource visualizer
- Settings
- Federated credentials**
- Properties
- Locks

Federated credentials

Configure an identity from an external OpenID Connect Provider to get tokens as this managed identity to access Microsoft Entra ID protected services. [Learn more about how to create an identity from an external OpenID](#)

+ Add Credential

3 of 20 configured

Name ↑	Issuer
azure-aks	https://eastus2.oic.prod-aks.azure.com/tenant-id/cluster-id

Configure an identity from an external OpenID Connect Provider to get tokens as this managed identity to access Microsoft Entra ID protected services.

Federated credential scenario \*

Configure a Kubernetes service account to get tokens as this application and access Azure resou...

[Configuration guide for Kubernetes identities](#)

## Connect your Kubernetes cluster

Please enter the details of the Kubernetes cluster that you want to connect to Microsoft Entra ID. These values will be used by Microsoft Entra ID to validate the connection and should match your Kubernetes OIDC configuration.

Cluster Issuer URL \*

https://eastus2.oic.prod-aks.azure.com/tenant-id/cluster-id

Namespace \*

workload-identity

Service Account \*

nymeria

Subject identifier

system:serviceaccount:workload-identity:nymeria

This value is generated based on the Kubernetes account details provided. [Edit \(optional\)](#)

## Credential details

Enter and review the details for this credential. The credential name cannot be edited after creation.

Name \*

azure-aks

Audience \*

api://AzureADTokenExchange

[Edit \(optional\)](#)

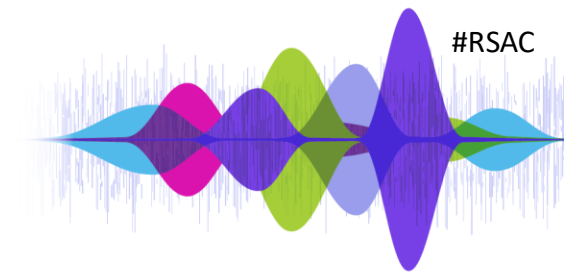


# Kubernetes Cross-Cloud Workload Identity Federation

Securing Cloud Access with Kubernetes  
Workload Identity

Many Voices.  
**One Community.**

# Default Identity Token Configuration



The default identity created when enabling workload identity for a cluster are not designed for federating into external APIs:

AWS Elastic Kubernetes  
Service (EKS)



Google Kubernetes  
Engine (GKE)



Azure Kubernetes  
Service (AKS)



Default Audience	Default Expiration
sts.amazonaws.com	1 Day *
Cluster Endpoint (self)	1 Year
api://AzureADTokenExchange	1 Hour *

\* Configurable using the workload identity annotations

# Kubernetes Google Projected Volume

The Google API volume mount configuration for pods running in EKS and AKS:

- Support for *secret*, *configMap*, *serviceAccountToken* data sources
- The *serviceAccountToken* source supports configuration options for setting the token's:
  - Audience (sts.googleapis.com)
  - Expiration (3600 seconds)
- Projected volume is mounted into the containers file system

```
1 volumes:
2   - name: gcp-token
3     projected:
4       sources:
5         - serviceAccountToken:
6           path: token
7           expirationSeconds: 3600
8           audience: "sts.googleapis.com"
9         - configMap:
10          name: gcloud-config
11          items:
12            - key: config.json
13              path: config.json
14 containers:
15   ...
16   volumeMounts:
17     - name: gcp-token
18       mountPath: "/var/run/secrets/gcp/
19         serviceaccount"
```

# Google Cloud External Cluster Workload Identity

1. External Kubernetes cluster OIDC providers need to be added to the project's workload identity pool with the issuer and audience:



Filter Enter property name or value

Display Name	ID	Providers	Status
nymeria-k8s-383e58d5	nymeria-k8s-383e58d5	2	✓
AWS EKS cluster	aws-eks-383e58d5	OIDC	✓
Azure AKS cluster	azure-aks-383e58d5	OIDC	✓

**Provider details**

Name \*  
AWS EKS cluster

ID  
aws-eks-383e58d5

Issuer (URL) \*  
https://oidc.eks.us-west-1.amazonaws.com/id/cluster-id

Issuer URL must start with https://

2. External identities federating through the workload identity provider can be used in direct role bindings:

```

1 bindings:
2   role: roles/storage.objectViewer
3   - members:
4     - principal://iam.googleapis.com/projects/#####/locations/global/
5       workloadIdentityPools/nymeria-k8s/
6       subject/system:serviceaccount:workload-identity:nymeria
  
```

# Kubernetes AWS Projected Volume


The AWS API volume mount configuration for pods running in AKS and GKE:

- The *serviceAccountToken* source configures the identity token:
  - Audience (sts.amazonaws.com)
  - Expiration (3600 seconds)
- Projected volume is mounted into the containers file system
- Pod environment variables are set to match the AWS IRSA workload identity configuration

```
1  volumes:
2    - name: aws-token
3      projected:
4        sources:
5          - serviceAccountToken:
6              path: token
7              expirationSeconds: 3600
8              audience: "sts.amazonaws.com"
9  containers:
10    ...
11    volumeMounts:
12      - name: aws-token
13        mountPath: "/var/run/secrets/aws/
14          serviceaccount"
15    env:
16      - name: AWS_ROLE_ARN
17        value: arn:aws:iam::0123456789012:
18              role/nymeria
19      - name: AWS_WEB_IDENTITY_TOKEN_FILE
20        value: /var/run/secrets/aws/
21              serviceaccount/token
```

# AWS Cloud External Cluster Workload Identity

1. External Kubernetes cluster OIDC providers need to be added to the AWS account's identity providers with the issuer and audience:



Provider	Type
<a href="https://container.googleapis.com/v1/projects/my-project-id/locations/us-west2/clusters/nyme">container.googleapis.com/v1/projects/my-project-id/locations/us-west2/clusters/nyme</a>	OpenID Connect
<a href="https://eastus2.oic.prod-aks.azure.com/tenant-id/cluster-id/">eastus2.oic.prod-aks.azure.com/tenant-id/cluster-id/</a>	OpenID Connect
<input type="radio"/> <a href="https://oidc.eks.us-west-1.amazonaws.com/id/my-cluster-id">oidc.eks.us-west-1.amazonaws.com/id/my-cluster-id</a>	OpenID Connect

2. Add the external Kubernetes cluster identity provider principals to the IAM Role trust policy to grant external pod service account permissions to assume the role:

```

1  { "Effect": "Allow",
2    "Action": "sts:AssumeRoleWithWebIdentity",
3    "Principal": {
4      "Federated":
5        "arn:aws:iam::123456789012:oidc-provider/eastus2.oic.prod-
6          aks.azure.com/tenant-id/cluster-id/" },
7    ...

```



# Kubernetes Azure Projected Volume

The Azure API volume mount configuration for pods running in EKS and GKE:

- The *serviceAccountToken* source configures the identity token:
  - Audience  
(api://AzureADTokenExchange)
  - Expiration (3600 seconds)
- Projected volume is mounted into the containers file system
- Pod environment variables are set to match the Azure workload identity configuration

```
1 volumes:
2   - name: azure-token
3     projected:
4       sources:
5         - serviceAccountToken:
6           path: token
7           expirationSeconds: 3600
8           audience:
9             "api://AzureADTokenExchange"
10 containers:
11   ...
12   volumeMounts:
13     - name: azure-token
14       mountPath: "/var/run/secrets/azure/
15         serviceaccount"
16   env:
17     - name: ARM_TENANT_ID ...
18     - name: ARM_CLIENT_ID ...
19     - name: AZURE_FEDERATED_TOKEN_FILE
20       value: "/var/run/secrets/azure/
21         serviceaccount/token
```

# Azure Cloud External Cluster Workload Identity

External Kubernetes cluster OIDC providers need to be added to the Azure managed identity federated credentials with the issuer and audience:

**nymeria-49d0fec8 | Federated credentials** ☆ ...

Managed Identity

Search

Overview

Activity log

Access control (IAM)

Tags

Azure role assignments

Associated resources (preview)

Resource visualizer

Settings

**Federated credentials**

Properties

Locks

**Federated credentials**

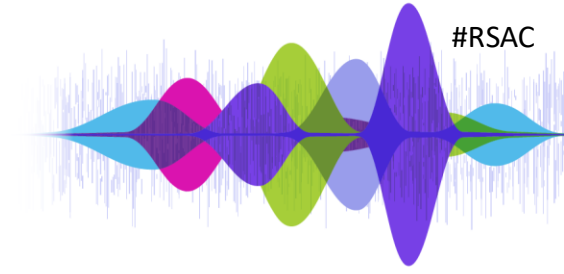
Configure an identity from an external OpenID Connect Provider to get tokens as this managed identity to access Microsoft Entra ID protected services. [Learn more about how to create an identity from an external OpenID](#)

+ Add Credential

3 of 20 configured

Name ↑	Issuer	Subject Identifier	Delete
<a href="#">gcp-gke</a>	https://container.googleapis.com...	system:serviceaccount:workload-...	
<a href="#">aws-eks</a>	https://oidc.eks.us-west-1.amazo...	system:serviceaccount:workload-...	
<a href="#">azure-aks</a>	https://eastus2.oic.prod-aks.azur...	system:serviceaccount:workload-...	

# Apply What You Have Learned Today



- Next week you should:
  - Clone the Nymeria repository and review the documentation
  - Scan your Kubernetes clusters for static credentials stored in *Secrets* and *ConfigMaps*
- Within three months, you should:
  - Configure Kubernetes pods to use intra-cloud integrated workload identity capabilities where possible
- Within six months, you should:
  - Expand Kubernetes workload identity using cross-cloud projected volumes to eliminate cross-cloud static credentials

**RSAC** | 2025  
Conference

Many Voices.  
**One Community.**

**Thank you for attending!**

**Eric Johnson**  
Principal Security Engineer, Puma Security  
Senior Instructor, SANS Institute  
[linkedin.com/in/eric-m-johnson/](https://www.linkedin.com/in/eric-m-johnson/)